

how your **AWS SageMaker** model is accessed and how authentication works.

1. Inside AWS

Your flow is:

Model → Endpoint → Prediction

The **model** is your trained insurance prediction model.

The **endpoint** is the AWS service that makes your model available for prediction requests.

You have two endpoint deployment types:

- **Server-based inference** → a server/instance runs continuously. You generally pay while it is running.
- **Serverless inference** → AWS runs compute when requests arrive, instead of keeping a traditional inference instance running.

So:

Model → Endpoint → Server-based / Serverless

2. AWS dashboard / SageMaker

Your drawing also indicates that inside AWS you manage things through:

AWS Dashboard → SageMaker

From SageMaker, you can manage:

- Models
- Endpoints
- Endpoint configurations
- Studio/Jupyter environments
- Other SageMaker resources

3. Outside AWS — IAM

The bottom part is about **IAM authentication**.

The idea is:

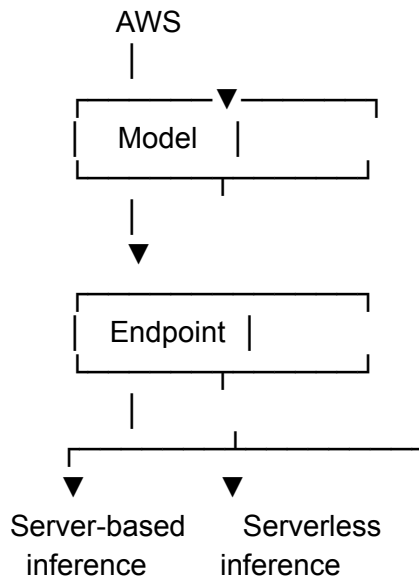
IAM User → Access Keys → AWS services

An IAM user can have an **Access Key ID** and **Secret Access Key**.

These keys can be used by applications such as Python/Boto3 to authenticate with AWS.

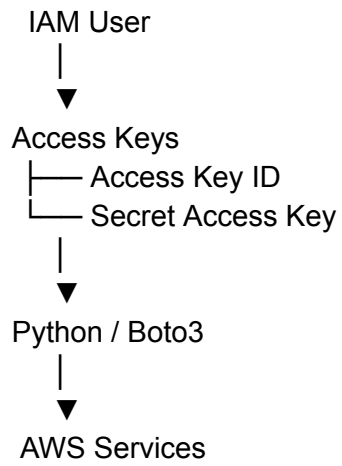
⚠ Never share the Secret Access Key publicly. If someone gets it, they may potentially access AWS resources according to the permissions attached to that IAM user.

In one simple diagram



=====

Outside / Authentication



model deployment (SageMaker Endpoint) from AWS authentication (IAM Access Keys)

=====